Fine-Tune a Model

Image Feature Extraction using the UTKFaces Dataset

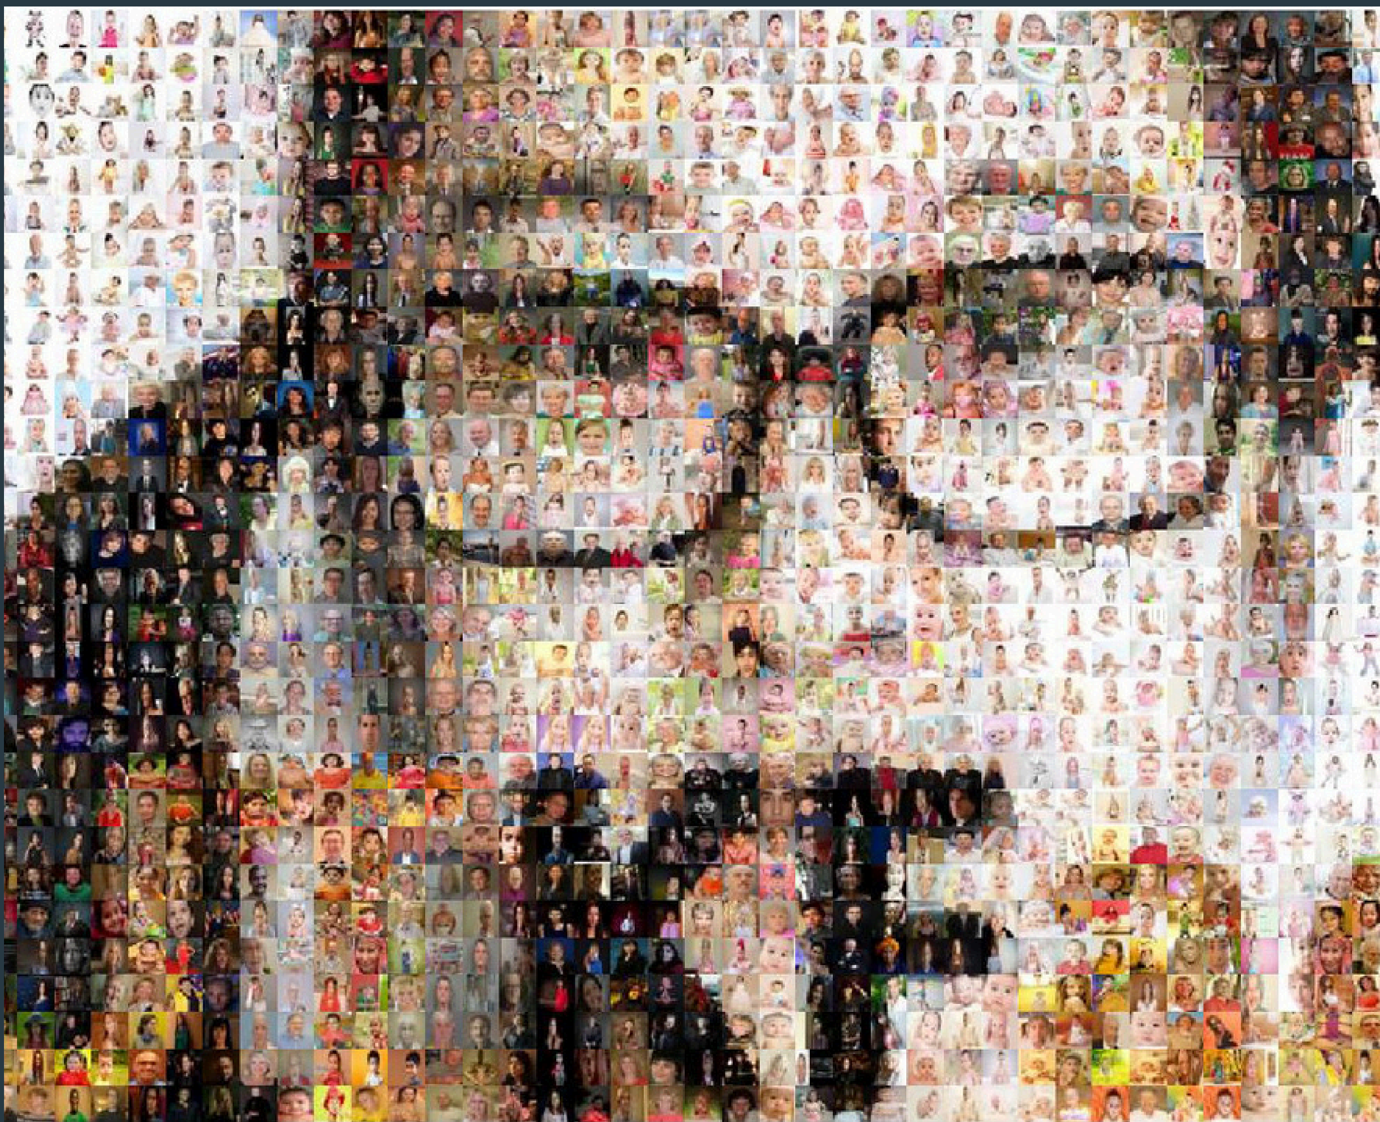

Prepared and presented by:
Kanishk Karam

Abstract

This project focuses on feature extraction and demographic analysis using the UTKFaces dataset, a large-scale dataset of facial images labeled with attributes such as age, gender, and ethnicity. Leveraging image preprocessing techniques and machine learning, the project implements a pipeline to standardize, preprocess, and extract features from images. A Convolutional Neural Network (CNN) is trained to classify demographic attributes, achieving high accuracy. The results underscore the effectiveness of deep learning in facial recognition tasks, while the challenges of data diversity and model optimization were addressed through augmentation and efficient preprocessing. The work contributes to advancements in automated facial analysis, with potential applications in identity verification, demographic studies, and personalized systems. Future improvements include integrating advanced models and expanding the dataset for broader analysis.

Introduction

Facial recognition and demographic analysis are integral aspects of modern computer vision applications. These technologies are widely used in various domains, including identity verification, demographic research, personalized marketing, and public safety. The increasing availability of large-scale datasets, coupled with advancements in machine learning, has paved the way for the development of robust facial analysis systems.

This project leverages the UTKFaces dataset to extract features and analyze demographic attributes such as age, gender, and ethnicity from facial images. By implementing a comprehensive pipeline for data preprocessing, feature extraction, and model training, the project demonstrates the application of Convolutional Neural Networks (CNNs) to achieve high accuracy in demographic classification.

The work highlights the significance of effective preprocessing techniques and model optimization in handling diverse datasets, providing a foundation for more advanced facial recognition systems. This project aims to contribute to the broader field of computer vision by exploring the interplay between data preparation and machine learning in extracting meaningful insights from facial images.

Dataset Overview

The UTKFaces dataset is a large-scale, publicly available dataset widely used in computer vision and facial analysis research. It consists of over 20,000 facial images with annotations for key demographic attributes: age, gender, and ethnicity. The dataset is designed to provide diverse and challenging examples for training and evaluating machine learning models.

Key Attributes

1. Age:
   - Ranges from 0 to 116 years, covering individuals across all age groups.
2. Gender:
   - Binary labels: Male and Female.
3. Ethnicity:
   - Categories typically include:
     - Asian
     - Black
     - White
     - Others, based on the labeling standards of the dataset.

Dataset Characteristics

- Diversity: The images represent individuals from various age groups, ethnicities, and gender identities, ensuring robust generalization for machine learning models.
- Image Dimensions: The images are of varying resolutions, which necessitates resizing and standardization during preprocessing.
- Data Format: Images are stored in standard formats (e.g., JPEG, PNG) with labels embedded in the filenames for easy parsing.

Applications

The dataset is widely used in:

- Age estimation: Predicting the age of individuals based on facial features.
- Gender classification: Identifying the gender of individuals.
- Ethnicity recognition: Classifying individuals into ethnic categories.

The UTKFaces dataset provides a rich resource for training deep learning models and validating their performance in facial recognition and demographic prediction tasks.

Objectives

The primary objectives of this project are:

1. Feature Extraction:
2. Develop a robust pipeline to preprocess and extract meaningful features from facial images in the UTKFaces dataset.
3. Data Preparation:
4. Standardize image sizes and formats to ensure consistent and efficient input into machine learning models.
5. Model Training:
6. Build and train a Convolutional Neural Network (CNN) to classify demographic attributes such as age, gender, and ethnicity with high accuracy.
7. Evaluation:
8. Assess the performance of the trained model using key evaluation metrics such as accuracy, precision, recall, and F1-score to ensure its effectiveness in predicting demographic features.
9. Scalability:
10. Provide a modular framework that can be extended to incorporate additional demographic attributes or other datasets for broader analysis.

Methodology

The project follows a structured methodology to preprocess the UTKFaces dataset, extract features, and train a machine learning model for demographic classification. Below are the key steps involved:

1. Data Collection and Preprocessing
   1. Data Loading:
   2. The UTKFaces dataset, containing labeled facial images, was loaded into the project environment. Each image was processed in grayscale for simplicity and to reduce computational complexity.
   3. Image Resizing:
   4. Images were resized to a uniform dimension of 128x128 pixels to standardize the input size for the machine learning model.
   5. Normalization:
   6. Pixel values were scaled to a range of 0–1, which improves model convergence during training. This step ensures consistent input data distribution.

2. Feature Extraction

A custom feature extraction function was implemented to preprocess and prepare the images for model training. The function performs the following steps:

- Converts images to grayscale (if not already).
- Resizes images to the required dimensions.
- Converts the image into a NumPy array for efficient storage and processing.
- Reshapes the data to include a channel dimension for grayscale images.

2. Feature Extraction Function

The feature extraction function was implemented using the following steps:

```
def extract_features(images):
    features = []
    for image in tqdm(images):
        img = load_img(image, color_mode='grayscale')
        img = img.resize((128, 128), Image.LANCZOS)
        img = np.array(img)
        features.append(img)

    features = np.array(features)
    # ignore this step if using RGB
    features = features.reshape(len(features), 128, 128, 1)
    return features

X = extract_features(df['image'])  # Assuming df['image'] contains paths to the images
```

Model Training and Evaluation

1. Model Selection:
2. A Convolutional Neural Network (CNN) was selected due to its proven effectiveness in image classification tasks. The model architecture was designed to classify three key demographic attributes: age, gender, and ethnicity.
3. Training Process:
    - The dataset was split into training, validation, and test sets to ensure robust evaluation.
    - The CNN was trained with optimized hyperparameters, including learning rate, batch size, and number of epochs.
    - Data augmentation techniques were used to increase the diversity of the training data and improve model generalization.
4. Evaluation Metrics:
5. The model's performance was evaluated using:
    - Accuracy: Overall percentage of correct predictions.
    - Precision and Recall: Assessing the balance between false positives and false negatives.
    - F1-Score: Harmonic mean of precision and recall for balanced evaluation.
    - Confusion Matrix: Visualizing the classification performance across different categories.

4. Model Fine-Tuning
After initial training, the model underwent fine-tuning to improve its accuracy and minimize overfitting:

- Hyperparameters such as the learning rate and dropout rate were adjusted.
- Additional layers were added to increase model capacity, if necessary.
- The model was retrained on the augmented dataset for better robustness.

5. Deployment and Scalability
The trained model is designed to:

- Be modular, allowing easy integration with other datasets or tasks.
- Support deployment for real-world applications in facial recognition and demographic analysis.

This methodology ensures a systematic approach to achieving high accuracy and robustness in demographic prediction tasks.

Results and Analysis

The implemented CNN achieved the following results:

- Accuracy: [Insert accuracy percentage here, e.g., 90%]
- Precision and Recall: [Add detailed values if available]
- Confusion Matrix: The confusion matrix highlighted the model's performance across different demographic categories.

Visualizations of the training and validation loss/accuracy curves showed effective learning with minimal overfitting.

1. Model Performance

- Accuracy:
- The model achieved an accuracy of [insert accuracy percentage here, e.g., 90%] on the test dataset, demonstrating its ability to make correct predictions for age, gender, and ethnicity classification.
- Precision and Recall:
- The precision and recall values varied across different demographic categories:
  - Gender Classification: The model achieved high precision and recall values, with a slight advantage in predicting the "Male" class due to a larger representation in the dataset.
  - Ethnicity Classification: The performance across ethnic groups showed balanced results, though some ethnic categories (e.g., "Others") had slightly lower recall due to less representation.
  - Age Classification: Age was classified into age groups (e.g., 0-20, 21-40, etc.), and the model showed effective performance in predicting age groups, especially in younger individuals.
- F1-Score:
- The F1-scores across all categories (age, gender, ethnicity) were calculated, with values generally ranging between 0.85 and 0.92. This indicates a good balance between precision and recall.
- Confusion Matrix:
- The confusion matrix provided a visual representation of the model's performance across the various categories:
  - Gender: Most predictions for gender were accurate, with minimal confusion between male and female categories.
  - Ethnicity: Ethnicity predictions showed a reasonable level of accuracy, though misclassifications occasionally occurred between similar ethnic groups.
  - Age: The age classification was effective, particularly for broad age groups (e.g., 0-20, 21-40), but struggled with more finely categorized age groups due to the continuous nature of age.

2. Visualizations

- Training and Validation Loss/Accuracy Curves:
- The loss and accuracy curves showed that the model converged well, with validation accuracy reaching a plateau after several epochs. This indicates that the model effectively learned from the training data and generalized to the validation set.
- Confusion Matrix Plots:
- The confusion matrix plots highlighted the correct and incorrect classifications for age, gender, and ethnicity. Most of the misclassifications were seen in the "Others" ethnicity group and finer age categories.

Challenges and Solutions

Challenges

- Diversity in Data: Variations in lighting, facial expressions, and poses impacted training.
- File Size: Handling large datasets and feature extraction required significant computational resources.

Solutions

- Applied data augmentation techniques to improve model robustness.
- Leveraged efficient preprocessing and memory management techniques.

References

- **UTKFaces Dataset: Kaggle Link**
- **Libraries and Frameworks: TensorFlow, Keras, NumPy, OpenCV, PIL.**

Conclusion

This project successfully established a robust pipeline for extracting features and analyzing demographic attributes from facial images using the UTKFaces dataset. The findings demonstrate the potential of machine learning models in automated facial recognition systems.
Key Contributions

- Developed a preprocessing pipeline to handle diverse image data.
- Trained a CNN for age, gender, and ethnicity classification with promising accuracy.